Project 4: Reinforcement Learning

April 10th, 2024
TAs: Gabriele Tiboni (*gabriele.tiboni@polito.it*), Andrea Protopapa (*andrea.protopapa@polito.it*)

0   Introduction

The main goal of this project is to get familiar with the reinforcement learning paradigm in the context of robotic systems. It also aims to dive into the problem of learning a control policy for a robot in simulation using state-of-the-art RL algorithms, while introducing the challenges of using that policy in the real world. In particular, the student will learn about the *sim-to-real transfer* problem in robot learning literature, namely the task of learning policies in simulation through RL that can be directly transferred to real-world hardware, avoiding costly interactions with real setups and speeding up the training time. During this project, the student will be simulating the sim-to-real transfer task in a *sim-to-sim* scenario, where a discrepancy between *source* (training) and *target* (test) domains is manually injected. The student will implement *domain randomization* of dynamics parameters (e.g. masses or friction coefficients), a popular strategy to learn robust policies that transfer well to the target domain.

Specifically, the student must go through the following steps:

1. **Preliminaries**: Read the provided material to get familiar with the Reinforcement Learning framework, the sim-to-real transfer challenge and the common techniques to perform an efficient transfer from simulation to reality;
2. **Train your first RL agent**: Implement two basic Reinforcement Learning algorithms in literature (i.e., REINFORCE and Actor-Critic Policy Gradient algorithms) to train your first RL agent;
3. **Lower/upper bound baselines**: Implement an RL training pipeline via third-party APIs to state-of-the-art reinforcement learning algorithms such as PPO and SAC, providing the baselines for the subsequent phase;
4. **Uniform Domain Randomization**: Implement Uniform Domain Randomization (UDR) to learn robust policies in the source domain and limit the loss of performance during the sim-to-real transfer;
5. **Project extension**: Propose and implement a novel extension of the project;

For each of the aforementioned steps, a *soft deadline* is provided to guide the student along a potential schedule of the work throughout the project. This timeline is not to be seen as mandatory, but only as a suggestion of a possible weekly workload to ensure timely completion of the project.

Guiding questions, hints, and external references are also included to assist the student in comprehending the various topics independently.

Your submission will consist of (1) a **PDF report** (with paper-style) which carefully and systematically describes the entire work carried out during the project, and (2) the **code** implementation. In particular, the report should contain a brief introduction, a related work section, a methodological section, an experimental section with all the results and discussions, and a final brief conclusion. Remember to also describe your own project extension in the report.
Follow this **link** to open and create the template for the report using Overleaf.

1 Preliminaries

Soft deadline: 25/04/2024

Before starting, you're asked to take some time to familiarize yourself with the framework of Reinforcement Learning, the sim-to-real transfer challenge and SOTA strategies to overcome it. More in detail:

- Read sections 1.-1.4, 1.6, 3.-3.8, of [1] to understand the general Reinforcement Learning framework;
- Watch introductory video on Reinforcement learning by DeepMind video
- Read article on introduction to Reinforcement learning by OpenAI [part 1, part 2, part 3]
- Read sections 1.-1.3, 3.-3.4, of [2]
- Read sections 1., 2., 3. of [3]
- Read debate on the sim-to-real transfer paradigm [4]
- Read [5], [6], [9], blog post to understand domain randomization for sim-to-real transfer
- Read this set of slides by Josh Tobin and this article regarding domain randomization for both vision and dynamics properties

2 Train your first RL agent

Soft deadline: 06/05/2024

Train a simple RL agent on the gym Hopper environment. This environment comes with an easy-to-use python interface which controls the underlying physics engine — MuJoCo — to model the robot.
The hopper is a one-legged robot model whose task is to learn how to jump without falling, while achieving the highest possible horizontal speed.

Task 1 - The Gym Hopper environment

Check out the provided code template and start playing around with the underlying Hopper environment. Get familiar with the `test_random_policy.py` script, the python interface for

MuJoCo, the [gym documentaGym Documentationtion](#), and the hopper environment overall. Finally answer the questions below.

Guiding Questions

- What is the state space in the Hopper environment? Is it discrete or continuous?
- What is the action space in the Hopper environment? Is it discrete or continuous?
- What is the mass value of each link of the Hopper environment, in the source and target variants respectively?

Hints
- *If you need any help answering the above questions try looking at the [Mujoco documentation](#) or the [gym documentation](#).*
- *Bodies defined in the environment:* `env.sim.model.body_names`
- *Mass of all the corresponding bodies:* `env.sim.model.body_mass`
- *Number of degrees of freedom (DoFs) of the robot:* `env.sim.model.nv`
- *Number of DoFs for each body:* `env.sim.model.body_dofnum`
- *Number of actuators:* `env.sim.model.nu`
- *See other attributes [here](#)*

2.1   Implement basic RL training algorithms

Implement from scratch two basic policy gradient reinforcement learning algorithms to train a simple control policy for the Hopper environment. Use `agent.py` for implementing the reinforcement learning algorithm itself (for example, the agent and policy classes). These classes are used to implement the main training loop in the `train.py` file. In particular, follow the tasks below, and make sure to go through the provided external resources:

Task 2 - REINFORCE (Vanilla Policy Gradient)

REINFORCE is a policy optimization algorithm that directly adjusts the parameters of a stochastic policy to maximize expected cumulative rewards and it operates through gradient ascent as optimization algorithm, iteratively searching for optimal parameters that maximize the objective function.

Background material:

- See Section 13.3-13.4 in [1] (if you want to get into theoretical details go through 13.-13.2 as well)
- See "REINFORCE" in [blog post](#)

Implement the following algorithms:

- REINFORCE without baseline,
- REINFORCE with a constant baseline $b = 20$

and compare their results.

Guiding Questions

- Analyze the performance of the trained policies in terms of reward and time consumption.
- How would you choose a good value for the baseline?

- How does the baseline affect the training, and why?

Task 3 - Actor Critic

Actor-Critic methods combine the advantages of value-based and policy gradient reinforcement learning.

Background material::

- See Section 13.5 in [1]
- See "Actor-Critic" in [blog post](#)

Implement the Actor-Critic Policy Gradient algorithm:

- refer to `train.py` as a starting point. It is okay to look at publicly available code for reference, but it's likely easier and more helpful to understand how to implement the code by yourself.

Guiding Questions

- Analyze the performance of the trained policies in terms of reward and time consumption.
- Compare the results with the REINFORCE algorithm you have previously obtained, highlighting any notable differences in terms of learning stability and convergence speed.

3   Lower/upper bound baselines

Soft deadline: 17/05/2024

The student will simulate the sim-to-real transfer scenario in a simplified sim-to-sim setting, as no work takes place on an actual real robot. In particular, two custom *domains* have been created ad-hoc: policy training takes place in the *source* environment and the student will transfer and test the policy on the *target* environment — which technically represents the real world. To simulate the reality gap, the source domain Hopper has been generated by shifting the torso mass by 1kg with respect to the target domain.

3.1   Implement advanced RL training pipelines

In this section, you'll make use of a third-party library to train an agent with state-of-the-art reinforcement learning algorithms such as PPO (Proximal Policy Optimization) and SAC (Soft Actor-Critic). Stable-Baselines 3 is a state-of-the-art RL library offering efficient and stable implementations of various algorithms. Specifically, it provides robust implementations for PPO and SAC, making it an ideal choice for experimentation due to its ease of use and reliability.

Task 4 - Implement PPO and SAC

Follow the steps below, and make sure to go through the provided external resources:

1. Create a new script using the third-party library [stable-baselines3](#) (sb3) and train the Hopper agent with **one** algorithm of choice between PPO [8] and SAC [7].
   a. [openAI article on PPO](#)

b. [openAI article on SAC](#)
c. Explanation [video](#) on PPO, explanation [video](#) on SAC.
2. You may use the provided template in `train_sb3.py` as a starting point. It is okay to look at publicly available code for reference, but it's likely easier and more helpful to study the sb3 documentation and understand how to implement the code by yourself.

Hints*:*

- *While PPO and SAC are more complex to understand, they will lead to better convergence during policy training and are likely less sensitive to hyperparameters.*
- *For the evaluation phase, refer to the "Evaluation Helper" in the [documentation](#).*
- *Feel free to adjust any hyperparameter such as learning rate or others as presented in the documentation if needed.*
- *You can use [callbacks for additional functionality](#), i.e., saving checkpoints, implementing early stopping, or integrating with TensorBoard or WandB for visualization.*

3.2 Training and testing

Task 5 - Train and test your policies

Train two agents with your algorithm of choice, on the *source* and *target* domains respectively. Then, test each model and report its average return over 50 test episodes. In particular, report results for the following "training→test" configurations:

- source→source,
- source→target (**lower bound**),
- target→target (**upper bound**).

Test with different hyperparameters and report the best results found together with the parameters used. The results will be the upper bound and lower bound for the following Domain Randomization phase.

Guiding Questions

- Why do we expect lower performances from the "source→target" configuration w.r.t. the "target→target"?
- If higher performances can be reached by training on the target environment directly, what prevents us from doing so (in a sim-to-real setting)?

4 Uniform Domain Randomization

Soft deadline: 24/05/2024

Implement Uniform Domain Randomization (UDR) for the link masses of the Hopper robot.
In this setting, UDR refers to manually designing a uniform distribution over the three remaining masses in the *source* environment (considering that the torso mass is fixed at -1 kg w.r.t. the target one) and performing training with values that vary at each episode (sampled appropriately from the chosen distributions).
The underlying idea is to force the agent to maximize its reward and solve the task for a range of multiple environments at the same time, such that its learned behavior may be robust to dynamics variations.

Note that, since the choice of the distribution is a hyperparameter of the method, the student has to manually try different distributions in order to expect good results on the target environment.

Task 6 - Implement Domain Randomization

Train a UDR agent on the *source* environment with the same RL algorithm previously used. Later test the policy obtained on both the *source* and *target* environments.

Guiding Questions

- Is UDR able to overcome the unmodelled effect (shift of torso mass) and lead to more robust policies w.r.t. the naive "source→target" baseline in task 5?
- Can you think of limitations or downsides of UDR?

Hints:

- `env.sim.model.body_mass[i]` *controls the mass of the i-th body in the Hopper environment. In particular, the torso mass value is* `env.sim.model.body_mass[1]`, *the thigh mass value is* `env.sim.model.body_mass[2]`, *and so on.*
- *To check out all body names:* `env.sim.model.body_names`
- *Remember not to randomize the torso mass!*

5 Project extension

At this stage, you are expected to work on your own extension to the project. This may include the implementation of any idea of yours to further improve the sim-to-real transfer in our simple scenario and you can carry out novel analysis on particular aspects of the reinforcement learning pipeline, which are not necessarily obvious. All your ideas must be well motivated and technically sound when implemented. Keep in mind that, rather than requiring you to obtain actual improvements, this step is for you to go beyond the project guidelines and get a feeling of a research-like approach.

We list a few exemplary working directions below to provide reference of the expected effort. However, you are free to propose any idea of yours as an extension to the project.

- Implement one Adaptive Domain Randomization method (eg., SimOpt, DROID, DROPO)

- Implement one curriculum learning method for Domain Randomization (e.g., AutoDR, ActiveDR, DORAEMON)

- Study and investigate policy adaptivity in Domain Randomization (Markovian vs. memory-based policies):

  - Background material:

    i. 2018, "Policy Transfer with Strategy Optimization"

    ii. 2021, "Understanding Domain Randomization for Sim-to-real Transfer"

- Propose your own project extension

References

[1] "Reinforcement Learning: An introduction (Second Edition)" by Richard S. Sutton and Andrew G. Barto, PDF

[2] Kober, J., Bagnell, J. A., & Peters, J. (2013). "Reinforcement learning in robotics: A survey". The International Journal of Robotics Research, PDF

[3] Kormushev, P., Calinon, S., & Caldwell, D. G. (2013). "Reinforcement learning in robotics: Applications and real-world challenges", PDF

[4] Höfer, S., Bekris, K., Handa, A., Gamboa, J. C., Golemo, F., Mozifian, M., ... & White, M. (2020). "Perspectives on sim2real transfer for robotics: A summary of the R: SS 2020 workshop", PDF

[5] J. Tobin, R. Fong, A. Ray, J. Schneider, W. Zaremba, and P. Abbeel, "Domain Randomization for Transferring Deep Neural Networks from Simulation to the Real World." arXiv, Mar. 20, 2017. PDF

[6] Peng, X. B., Andrychowicz, M., Zaremba, W., & Abbeel, P. (2018, May). "Sim-to-real transfer of robotic control with dynamics randomization", PDF

[7] T. Haarnoja, A. Zhou, P. Abbeel, and S. Levine, "Soft Actor-Critic: Off-Policy Maximum Entropy Deep Reinforcement Learning with a Stochastic Actor.", PDF

[8] Schulman, J., Wolski, F., Dhariwal, P., Radford, A., & Klimov, O. (2017). "Proximal policy optimization algorithms", PDF

[9] Muratore, Fabio, et al. "Robot learning from randomized simulations: A review." Frontiers in Robotics and AI 9 (2022)., PDF